

CSCE 580: Introduction to AI

Lecture 16 Machine Learning – NN, DL

PROF. BIPLAV SRIVASTAVA, AI INSTITUTE

10TH OCT, 2024

Carolinian Creed: “I will practice personal and academic integrity.”

Credits: Copyrights of all material reused acknowledged

Organization of Lecture 16

- Introduction Segment
 - Recap of Lecture 16
- Main Segment
 - Neural Networks
 - Deep Learning
 - Trust Issues
 - Adversarial Attacks
- Concluding Segment
 - Course Project Discussion
 - About Next Lecture – Lecture 17
 - Ask me anything

Introduction Section

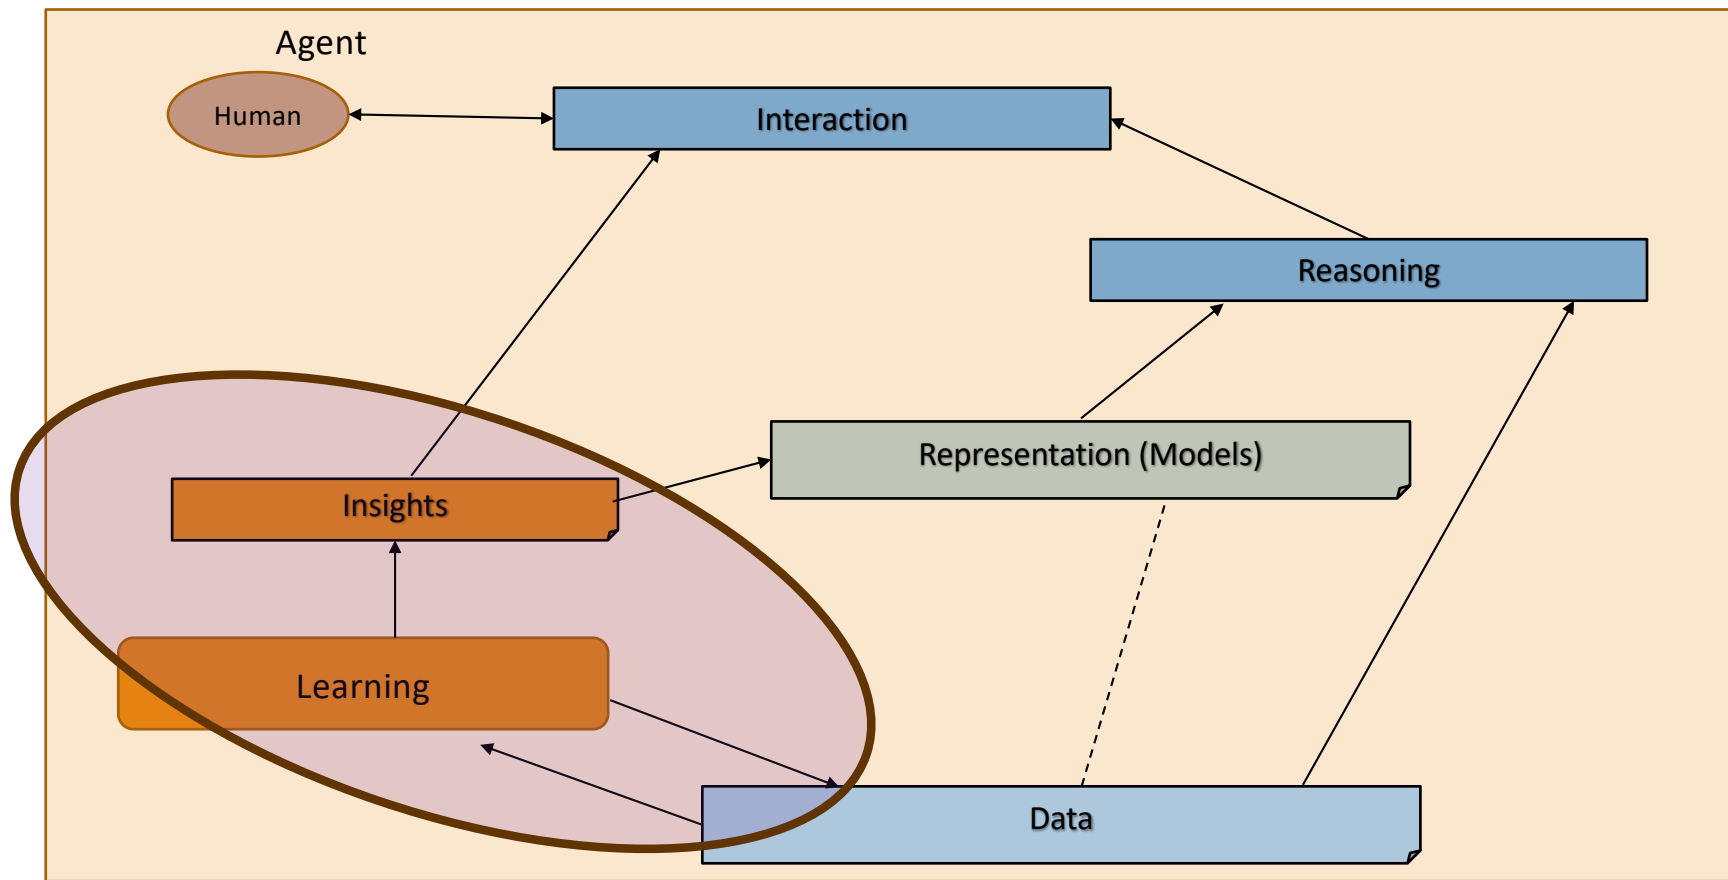
Recap of Lecture 18

- Topic discussed
 - Trust/ Explanations, LIME - Recap
 - Unsupervised ML Algorithms

Intelligent Agent Model



Relationship Between Main AI Topics



Where We Are in the Course

CSCE 580/ 581 – In This Course

- Week 1: Introduction, Aim: Chatbot / Intelligence Agent
- Weeks 2-3: Data: Formats, Representation and the Trust Problem
- Week 4-5: Search, Heuristics - Decision Making
- Week 6: Constraints, Optimization – Decision Making
- Week 7: Classical Machine Learning – Decision Making, Explanation
- Week 8: Machine Learning - Classification
- Week 9: Machine Learning - Classification – Trust Issues and Mitigation Methods
- Topic 10: Learning neural network, deep learning, Adversarial attacks
- Week 11: Large Language Models – Representation, Issues
- Topic 12: Markov Decision Processes, Hidden Markov models - Decision making
- Topic 13: Planning, Reinforcement Learning – Sequential decision making
- Week 14: AI for Real World: Tools, Emerging Standards and Laws; Safe AI/ Chatbots


Main Section

Credit: Retrieved from internet


Machine Learning – Insights from Data

- Descriptive analysis
 - Describe a past phenomenon
 - **Methods:** classification (feedback from label), clustering, dimensionality reduction, anomaly detection, neural methods, reinforcement learning (feedback from hint/ reward)
- Predictive analysis
 - Predict about a new situation
 - **Methods:** time-series, neural networks
- Prescriptive analysis
 - What an agent should do
 - **Methods:** simulation, reinforcement learning, reasoning
- New areas
 - Counterfactual analysis
 - Causal Inferencing
 - Scenario planning
 - Representation learning

Neural Network Methods



Model Depth and Learning Ability

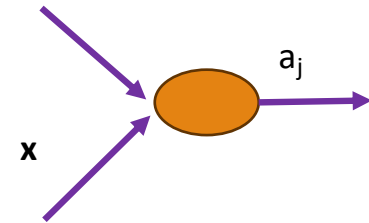


(a) A shallow model, such as linear regression, has short computation paths between inputs and output. (b) A decision list network has some long paths for some possible input values, but most paths are short. (c) A deep learning network has longer computation paths, allowing each variable to interact with all the others.

Adapted from:
Russell & Norvig, AI: A Modern Approach

Node (Unit) of a NN


- Notations and meanings
 - a_j : output of a unit j
 - $w_{i,j}$: weight of link from unit i to unit j
 - $a_j = g_j (\sum w_{i,j} a_i)$, where g_j is a nonlinear activation function
- $a_j = g_j (\mathbf{w}^T \mathbf{x})$, where \mathbf{w} is vector of weights leading into unit j and \mathbf{x} is the inputs to unit j




Popular Activation Functions

- Logistics or sigmoid function: $\sigma(x)$
 $= 1/(1 + e^{-x})$
- ReLU (rectified linear unit): $\max(0, x)$
- Softplus function: $\log(1 + e^x)$
 - Smooth version of ReLU
- $\tanh(x) = (e^{2x} - 1) / (e^{2x} + 1)$
 - Scaled and shifter version of sigmoid; $\tanh(x) = 2\sigma(2x) - 1$


a) the logistic or sigmoid function
b) the ReLU function and the softplus function
c) the tanh function.



(a)



(b)



(c)

Adapted from:
Russell & Norvig, AI: A Modern Approach

Note: All activation functions are non-linear

Loss functions

- Mean squared error

$$MSE = \frac{1}{n} \sum_{j=1}^n [f(X_j) - y_j]^2$$

- Categorical Cross Entropy


$$\text{Cost} = \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k [y_{ij} \log(\hat{y}_{ij})]$$

k is classes,
y = actual value
 \hat{Y} = prediction

- More loss functions:

<https://www.analyticsvidhya.com/blog/2022/06/understanding-loss-function-in-deep-learning/>

NN – Multi Layer Perceptron



Content and Image Courtesy:
<https://github.com/Thanasis1101/MLP-from-scratch>

(Stochastic) Gradient Descent

Gradient Descent

$\mathbf{w} \leftarrow$ any point in the parameter space

While not converged do:

For each w_i in \mathbf{w} do:

$$w_i \leftarrow w_i - \alpha \left(\frac{\partial}{\partial w_i} \text{Loss}(\mathbf{w}) \right)$$

Calculate the gradient of the loss function with respect to the weights along the gradient direction to reduce the loss.

Stochastic Gradient Descent (SGD)

Randomly select a small number of training examples at each step

Sources:

- https://en.wikipedia.org/wiki/Stochastic_gradient_descent
- Russell & Norvig, AI: A Modern Approach, Chapter 19

Logistic Regression in a Slide

Function estimate (linear)

W: weight, b: bias

$$f(X_j) = X_j W + b$$

Error Term (mean squared error)

$$MSE = \frac{1}{n} \sum_{j=1}^n [f(X_j) - y_j]^2$$

Update Weight

$$W^* = W - \eta \frac{dL}{dW}$$

Common Code Pattern

```
y = tf.matmul(x, W) + b
```

```
loss = tf.reduce_mean(tf.square(y - y_label))
```

NN Concepts

- **Epoch:** The number of times the learning algorithm will iterate over the entire dataset
- **Batch:** how many samples are processed before updating the model's internal parameters.
 - **Batch Gradient Descent:** Batch Size = Size of Training Set
 - **Stochastic Gradient Descent:** Batch Size = 1
 - **Mini-Batch Gradient Descent:** $1 < \text{Batch Size} < \text{Size of Training Set}$

Credit: <https://reentry.org/llm-training>

Universal Approximation Theorem

- A network with just two layers of computation units, first nonlinear, and the second linear, can approximate any continuous function to an arbitrary degree of accuracy.
- **Why:** a sufficiently large network can implement a lookup table for continuous functions
 - Nonlinear layer is the key

Sources:

- https://en.wikipedia.org/wiki/Universal_approximation_theorem
- Russell & Norvig, AI: A Modern Approach, Chapter 21

Datasets

- In keras, <https://keras.io/api/datasets/>
 - boston_housing
 - cifar10 module, cifar100, fashion_mnist, mnist
 - imdb module
 - reuters module
- In TF, https://www.tensorflow.org/datasets/catalog/overview#all_datasets

Keras Walkthrough

- Package: <https://keras.io/about/>
- Example model:
 - Sequential: https://keras.io/guides/sequential_model/
- Many examples: classification, image, text, audio
 - <https://keras.io/examples/>
- Future Keras: https://keras.io/keras_core/
 - Keras Core - run Keras workflows on top of TensorFlow, JAX, and PyTorch; preview of Keras 3.0

Code Examples With Keras and TF

1. Classification – diabetes

2. Try code

- Play with hyper-parameters
- Look at keras features used

Code location:

<https://github.com/biplav-s/course-ai-tai-f23/tree/main/sample-code/Class19-To-21-DL>

Discussion

- Impact of network structure:
 - Nodes / layer:
 - Layers:
 - Inter-connection structure:
- Impact of hyper-parameters:
 - Epochs:
 - Batch size:

Code Examples With Keras and TF

1. Classification – diabetes

2. Prediction/ representation learning – autoencoder

Code location:

<https://github.com/biplav-s/course-ai-tai-f23/tree/main/sample-code/Class19-To-21-DL>

Discussion

- Impact of network structure:
 - Nodes / layer:
 - Layers:
 - Inter-connection structure:
- Impact of hyper-parameters:
 - Epochs:
 - Batch size:

Code Examples With Keras and TF

1. Classification – diabetes
2. Prediction/ representation learning – autoencoder

3. Classification – MNIST

Code location:

<https://github.com/biplav-s/course-ai-tai-f23/tree/main/sample-code/Class19-To-21-DL>

Discussion

- Impact of network structure:
 - Nodes / layer:
 - Layers:
 - Inter-connection structure:
- Impact of hyper-parameters:
 - Epochs:
 - Batch size:

Keras and TensorFlow

- By Example:
 - <https://github.com/biplav-s/course-nl-f22/blob/main/sample-code/l11-nn-dl/Basic%20TensorFlow%20and%20Keras.ipynb>
 - TensorFlow's MNIST tutorial
 - <https://www.tensorflow.org/tutorials/quickstart/beginner>
- More examples
 - Number Addition by sequence learning: https://keras.io/examples/nlp/addition_rnn/
 - AutoEncoder: <https://machinelearningmastery.com/lstm-autoencoders/>

NN/ MLP

- Code examples:
 - <https://github.com/biplav-s/course-d2d-ai/blob/main/sample-code/l6-l7-l8-supervised-ml/Supervised-NaiveBayes-GradientBoost-NN-Classification.ipynb>
- Scikit Library:
 - MLP: https://scikit-learn.org/stable/modules/generated/sklearn.neural_network.MLPClassifier.html


Consideration: Which NN/DL Tool to Use

- See:
 - <https://www.simplilearn.com/keras-vs-tensorflow-vs-pytorch-article>
 - In theory, keras supports all major ones
 - Pytorch used in academic research more
 - TF used in production systems

Trust: Adversarial Attacks

Example (Gu et al. 2017)

- **ML Application:** Detect and classify street signs in images
- **Poisoning method:** Insert images where a special sticker is added to stop signs and the label changed to speed limit
- **Backdoor:** Adversaries ensure that any stop sign is misclassified simply by placing a sticker on it



Trust: Adversarial Attacks

- Cat and mouse on attacks and defenses
 - Example code: https://github.com/Trusted-AI/adversarial-robustness-toolbox/blob/main/notebooks/adversarial_training_mnist.ipynb
- Tools
 - Adversarial Robustness Toolbox (ART) - Python library for Machine Learning Security, <https://github.com/Trusted-AI/adversarial-robustness-toolbox>

Trust Issues with NN

- Robustness: can the model give the results in the presence of (input) perturbation? Noise?
- Computation/ footprint: why does the learning take so much compute resources?
- Data: is the data representative? How was the data obtained?
- Explainability: why does the model work?
- Fairness: Is the output fair to user groups?

Resources and Books

- Understanding Deep Learning, <https://udlbook.github.io/udlbook/>
- Deep Learning, Ian Goodfellow, Yoshua Bengio and Aaron Courville, <https://www.deeplearningbook.org/>
- AI – A Modern Approach, Russell & Norvig, <https://aima.cs.berkeley.edu/>
- Websites of libraries – Keras.

Course Project

Discussion: Projects

- New: two projects
 - Project 1: model assignment
 - Project 2: single problem/ llm based solving / fine-tuning/ presenting result

Lecture 19 & 20: Summary

- We talked about
 - Neural Networks
 - Deep Learning
 - Adversarial attacks
 - Trust Issues
- Others
 - Quiz 3 due today

Concluding Section

About Next Lecture – Lecture 17

Lecture 17: Text, Large Language Models

- Text processing
- Language Models (LMs)
- Large LMs

13	Oct 1 (Tu)	Machine Learning – Classification – Decision Trees, Random Forest, NBC, Gradient Boosting, ML-Text
14	Oct 3 (Th)	ML – Unsupervised / Clustering
15	Oct 8 (Tu)	Student presentations - project
16	Oct 10 (Th)	ML – NN, Deep Learning
17	Oct 15 (Tu)	Processing Natural Languages/ Language Models
	Oct 17 (Th)	
18	Oct 22 (Tu)	Large Language Models (LLMs) / Foundation Models
19	Oct 24 (Th)	Using LLMs – how and when ?